

# Docker and Security

Judith Kahrer

Some definitions

# What is a container?

*A container is a standard unit of software that packages up code and all its dependencies so the application runs quickly and reliably from one computing environment to another.*

[docker.com](https://docker.com)

*Container [...] is a method to package an application so it can be run, with its dependencies, isolated from other processes.*

[www.techradar.com](http://www.techradar.com)

*Containers provide a standard way to package your application's code, configurations, and dependencies into a single object.*

*Containers share an operating system installed on the server and run as resource-isolated processes, ensuring quick, reliable, and consistent deployments, regardless of environment.*

[aws.amazon.com](https://aws.amazon.com)

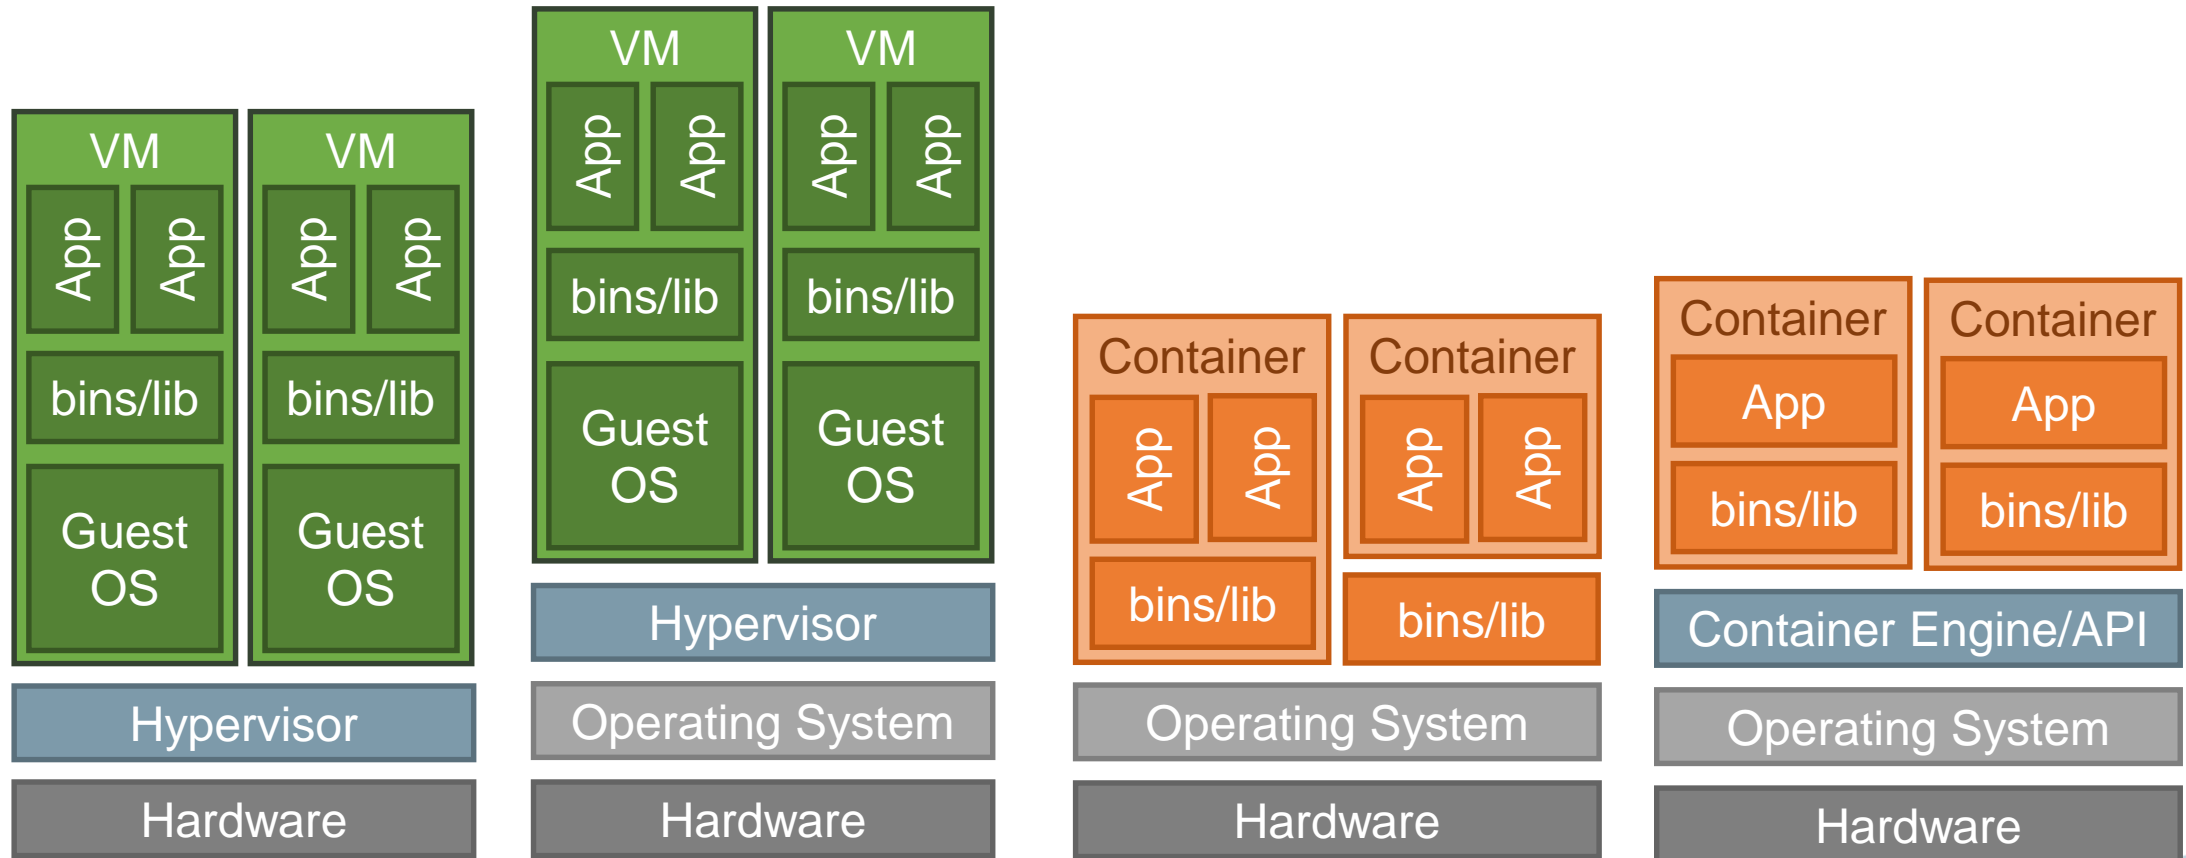
- Self-sufficient packages that enable reliable, consistent deployment of an application regardless of environment.
- Containers run as isolated processes on a target system.

# In technical terms

- *It's a lightweight VM – it 'feels' like it!*
- *It's chroot on steroids*

Jérôme Petazzoni, Docker Inc.


# Virtualization



[opennebula.org](http://opennebula.org) and [blog.netapp.com](http://blog.netapp.com)



# Timeline



[blogs.msdn.microsoft.com](https://blogs.msdn.microsoft.com)

Getting started

# Hello Docker

```
$ docker run hello-world
```

```
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
d1725b59e92d: Pull complete
Digest: sha256:0add3ace90ecb4adbf7777e9aacf18357296e799f81cabc9fde470971e499788
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.  
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

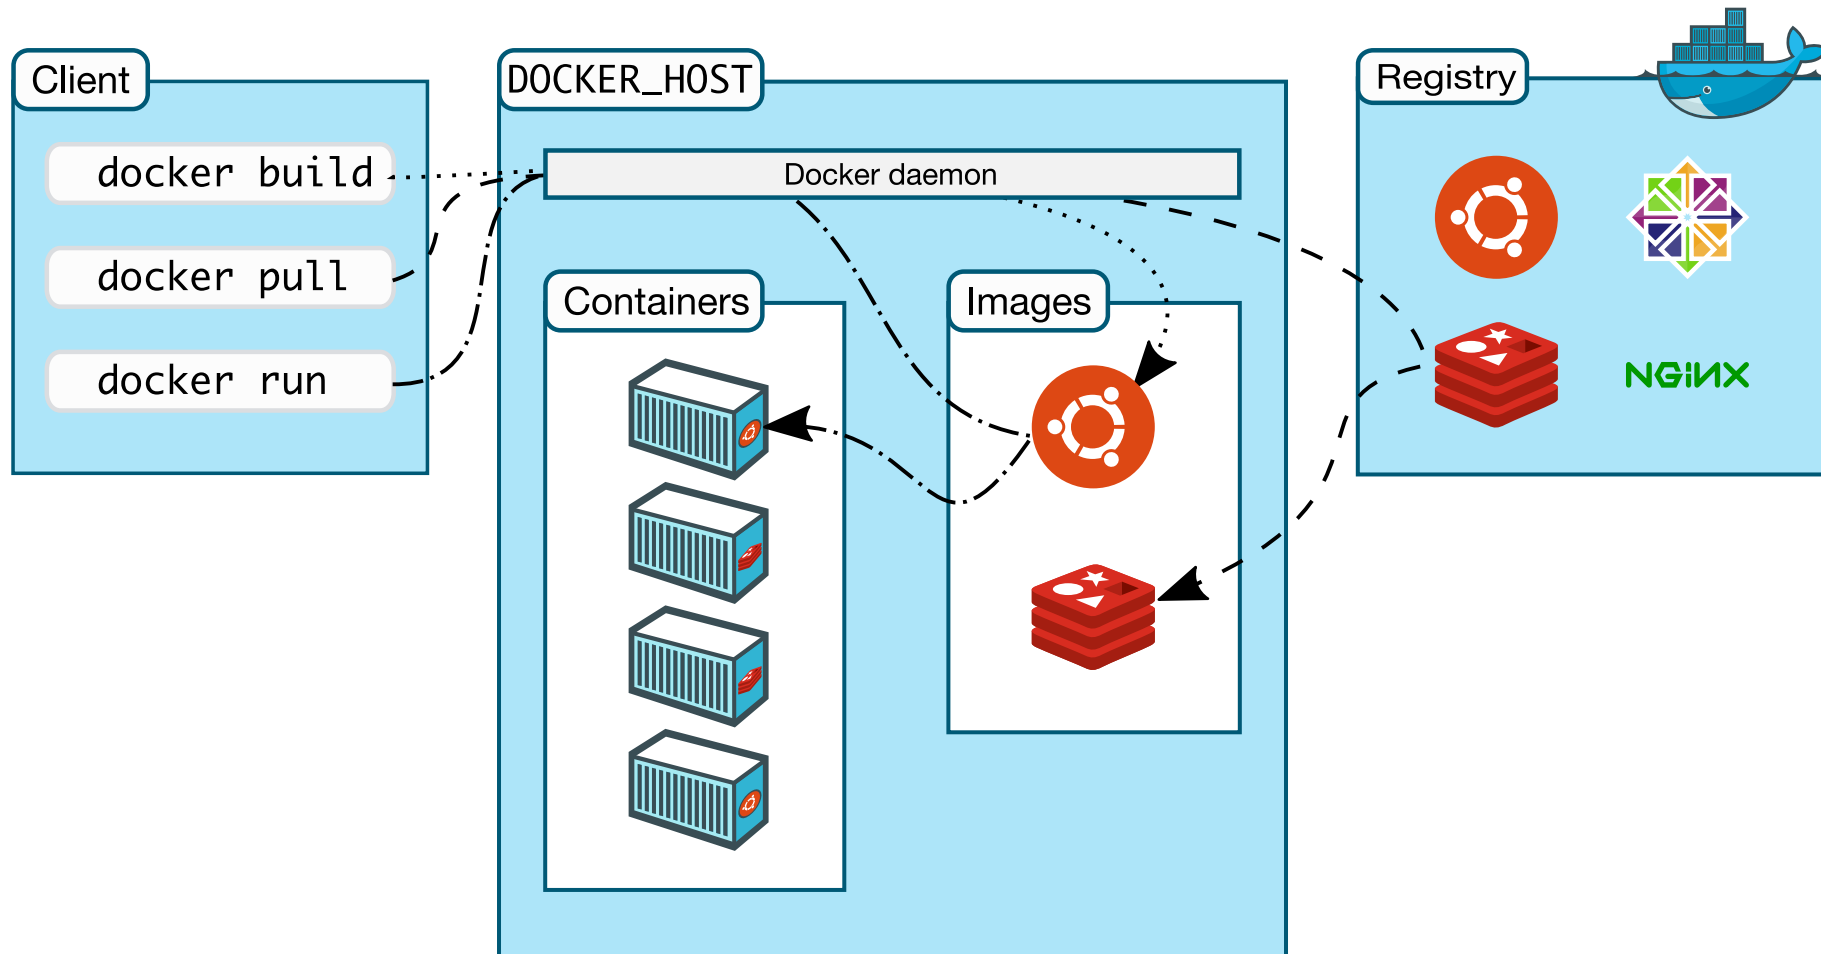
Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/get-started/>

# Architecture



# Docker Hub

- Public and private repositories
- Free-to-use registry
- Community Docker images

# Malicious Image

- Over 17 malicious Docker images
- Cryptomining containers
- Active for about 1 year
- 5 million pulls
- ~\$900000

[kromtech.com](http://kromtech.com) (june 2018)

# Image Security

- Official images
  - <https://hub.docker.com/official>
- Docker Content Trust
- Docker Trusted Registry and Notary
- Docker Security Scanning



Enabling

# Docker Content Trust

```
$ export DOCKER_CONTENT_TRUST=1
```

```
$ docker pull docker/trusttest
```

remote trust data does not exist [...]

```
$ docker pull busybox
```

Know the Basics

# Building Images

# Automated Builds

- Link Docker Hub with GitHub/Bitbucket
- Traceability between
  - Dockerfile
  - Version of the image

# Dockerfile

- Text file with instructions to build an image
- May contain i.a.
  - Base image
  - Environment variables
  - Files to include
  - Entrypoint and commands to run



FROM busybox

ENV variable=visible


COPY somelocalfile /greetings/fromMe.txt

CMD ["/bin/cat", "/greetings/fromMe.txt"]

# Layers

- Image Layers
  - Instructions from Dockerfile as read-only layers
  - Differences from previous layer
- Container Layer
  - Writeable layer on top during runtime

# Layers



[docs.docker.com](https://docs.docker.com)

```
$ docker history IMAGE
```

IMAGE	[...]	CREATED BY
11d8f9e12e42	[...]	CMD ["/bin/cat" "/greetings/fromMe.txt"]
63dc92745629	[...]	COPY file:3bfff68af9383544cc9f1cd820647b69...
09cdd4f30470	[...]	ENV variable=visible
788edf1f3e	[...]	CMD ["sh"]
<missing>	[...]	ADD file:63eebd629a5f7558c361be0305df5f16b...

Get into the Internals

# Kernel Features

# Inside Docker

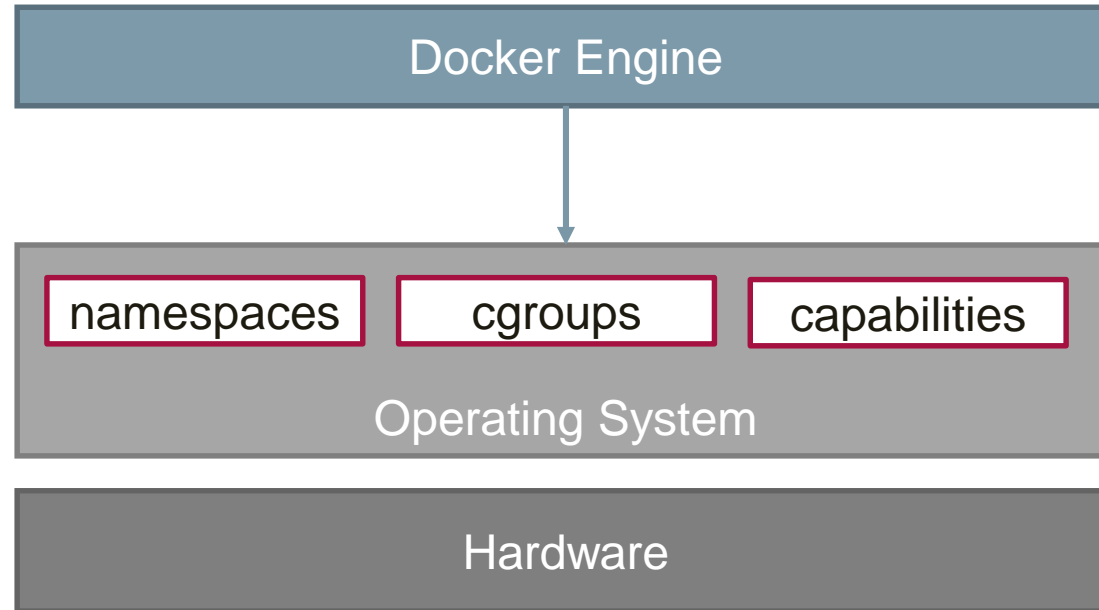
- Docker containers create an isolated complete execution environment

# Building Blocks

- Isolation
  - by namespaces
- Resource utilization
  - with the help of control groups (cgroups)
- Security
  - in form of capabilities



# Inside Docker

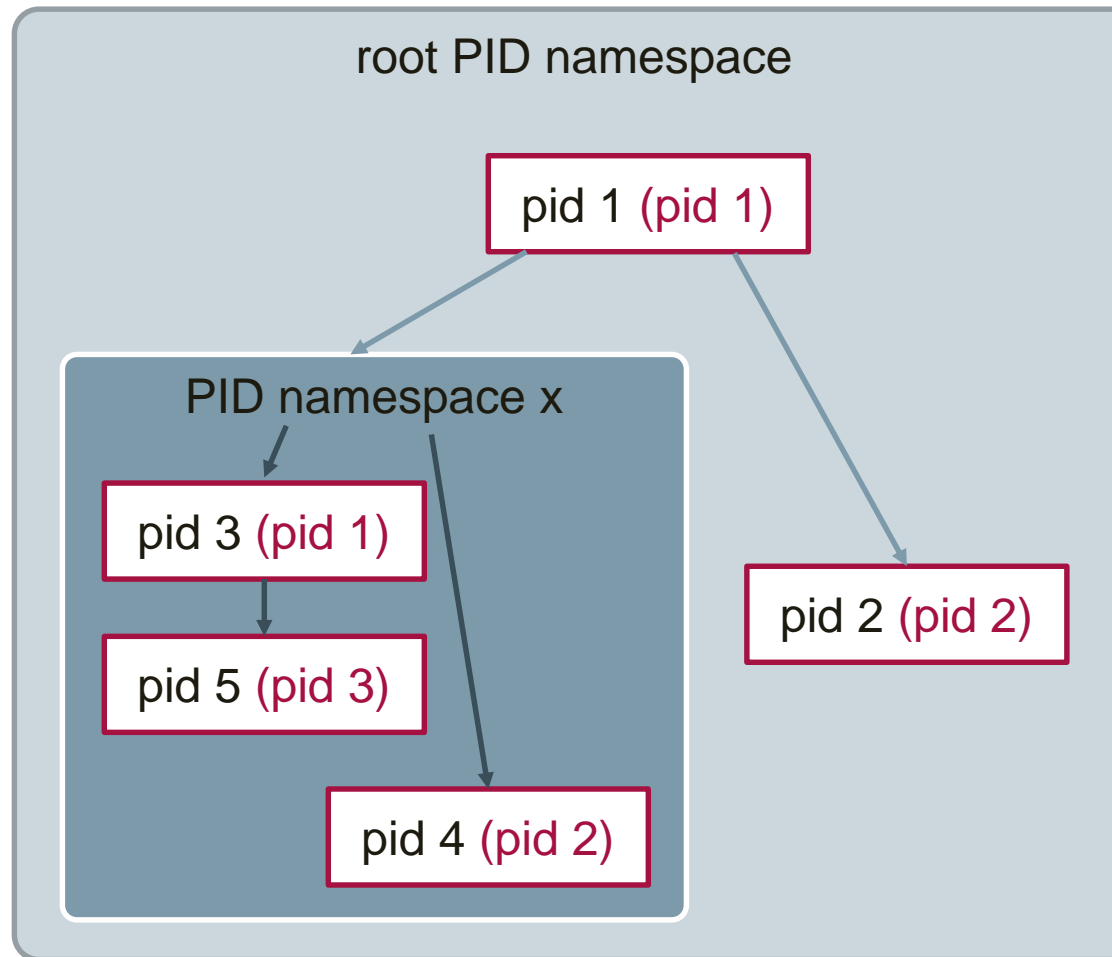


[medium.com](https://medium.com)

# Namespaces

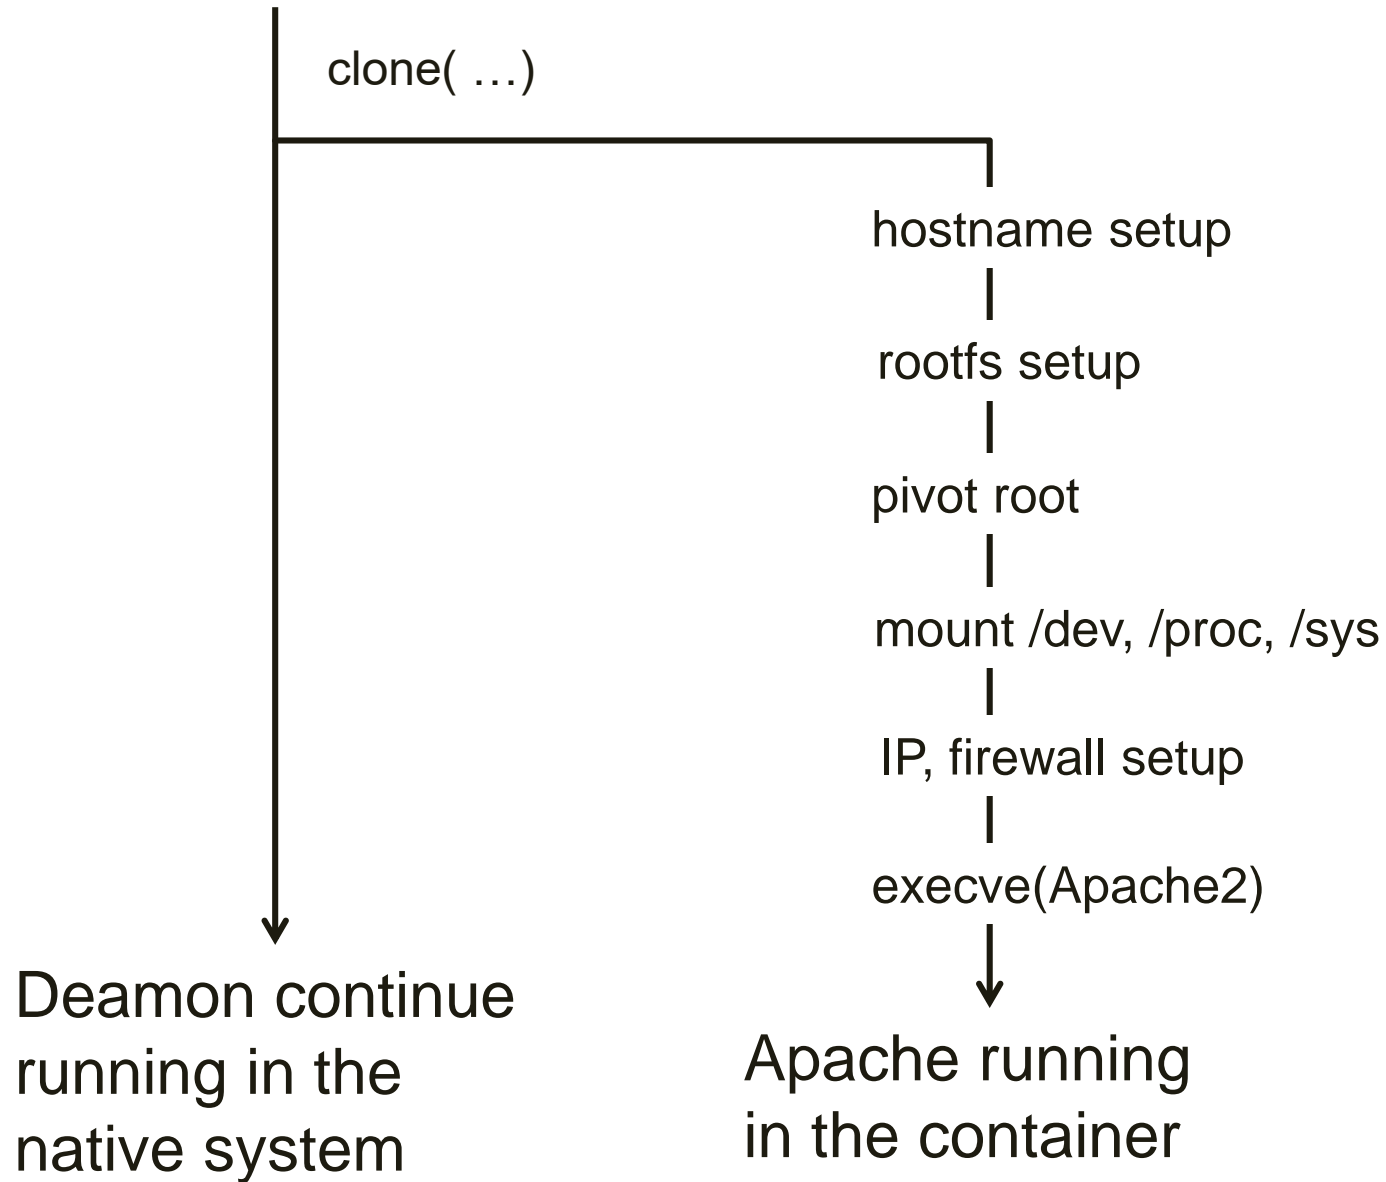
- Provide processes with their own view of the system
  - mnt (mount points, filesystems)
  - pid (processes)
  - uts (hostname)
  - ipc (interprocess communication)
  - net (network stack)
  - user (UIDs)

# PID namespace



black: real PID  
red: PID in namespace

# Docker Deamon



Security Namespace: Making  
Linux Security Frameworks  
Available to Containers, Sun et  
al. (2018)

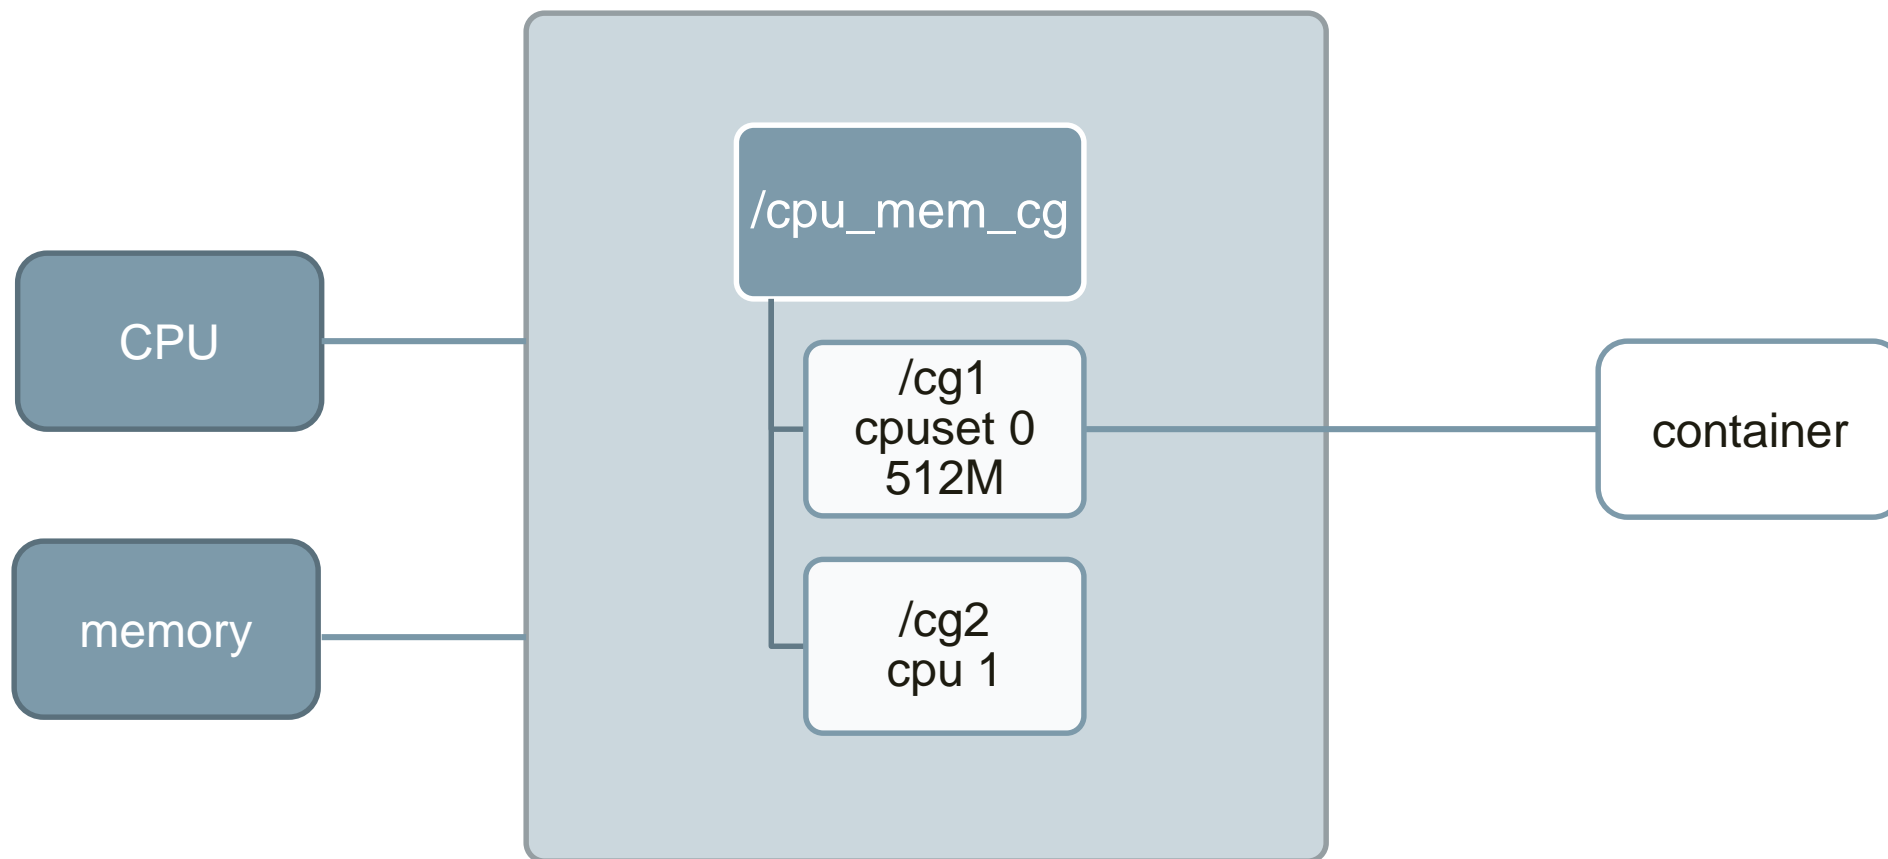
# cgroups

- Hierarchical groups of processes
- Resource accounting and limiting
  - memory
  - CPU time
  - disk I/O
  - network

# Limit Resources

- `--memory=<value>`
- `--memory-swap=<value>`
- `--oom-kill-disable`
- `--cpus=<value>`
- `--cpuset-cpus=<value>`
- ...

# cgroup example



[access.redhat.com](https://access.redhat.com)

```
$ docker run --cpuset-cpus 0 -m 512MB <id>
```



```
$ docker stats --no-stream
```

# Capabilities

- Docker container runs as root (default)
  - with a subset of capabilities
- Divide the power of superuser

# Default capabilities

- CHOWN
- DAC\_OVERRIDE
- FOWNER
- NET\_RAW
- ...

```
$ docker run --cap-drop ALL --user <uid>:<gid> <id>
```

# Secure Computing Mode

- seccomp
- Restrict system calls from the container
- More fine-grained control than capabilities
- Docker provides default whitelist

# Docker Awareness in Java

- CPU and memory config directly from the underlying host
  - Effects thread pool size
  - OutOfMemoryError
- 
- Mitigated in Java 8u131
  - Solved in Java 10

[efekahraman.github.io](https://efekahraman.github.io)

# DockerTest.java

```
public class DockerTest {  
    public static void main(String[] args) throws InterruptedException {  
        Runtime runtime = Runtime.getRuntime();  
        int cpus = runtime.availableProcessors();  
        long mmax = runtime.maxMemory() / 1024 / 1024;  
  
        System.out.println("System properties");  
        System.out.println("Cores          : " + cpus);  
        System.out.println("Memory (Max): " + mmax);  
        while (true) Thread.sleep(1000);  
    }  
}
```

```
$ docker run <id>
```

```
System properties Cores : 2
```

```
Memory (Max) : 241
```



```
$ docker run --cpu-shares 512 --memory 512MB  
<id>
```

```
System properties Cores : 2
```

```
Memory (Max) : 241
```

```
$ docker run \  
--cpuset-cpus 0 \  
--memory 512MB \  
--env JAVA_OPT="-XX:+UnlockExperimentalVMOptions \  
                -XX:+UseCGroupMemoryLimitForHeap" \  
<id>
```

```
System properties  
Cores          : 1  
Memory (Max) : 123
```

```
$ docker run --cpu-shares 512 --memory 512MB  
<id>
```

```
System properties Cores : 1
```

```
Memory (Max) : 123
```

```
$ docker run \  
--cpu-shares 512 \  
--memory 512MB \  
--env JAVA_OPT=-XX:-UseContainerSupport \  
<id>
```

System properties

Cores : 2

Memory (Max) : 241

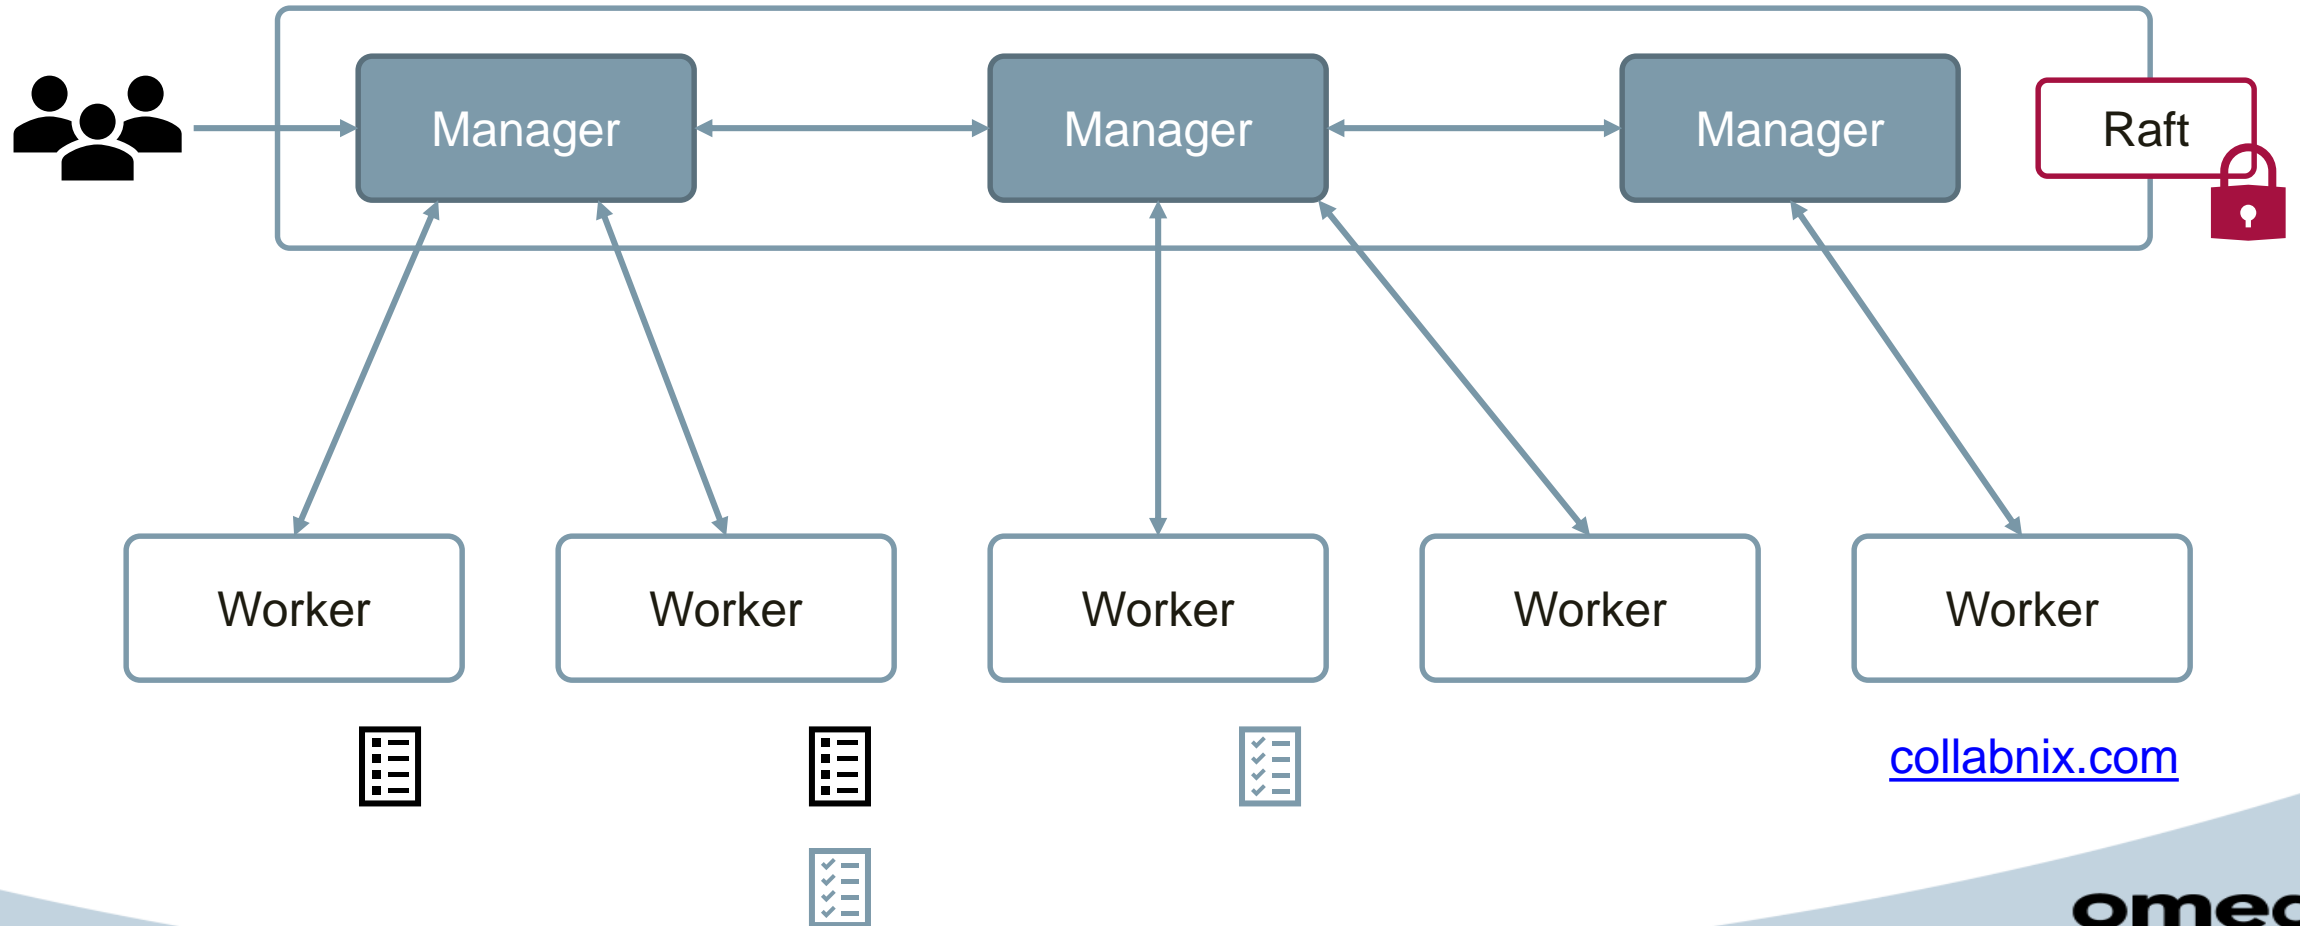
Orchestration

# Swarm mode

# Key Concepts

- Nodes
  - Manager
  - Worker
- Services
- Stack

# Cluster



# Security Concepts

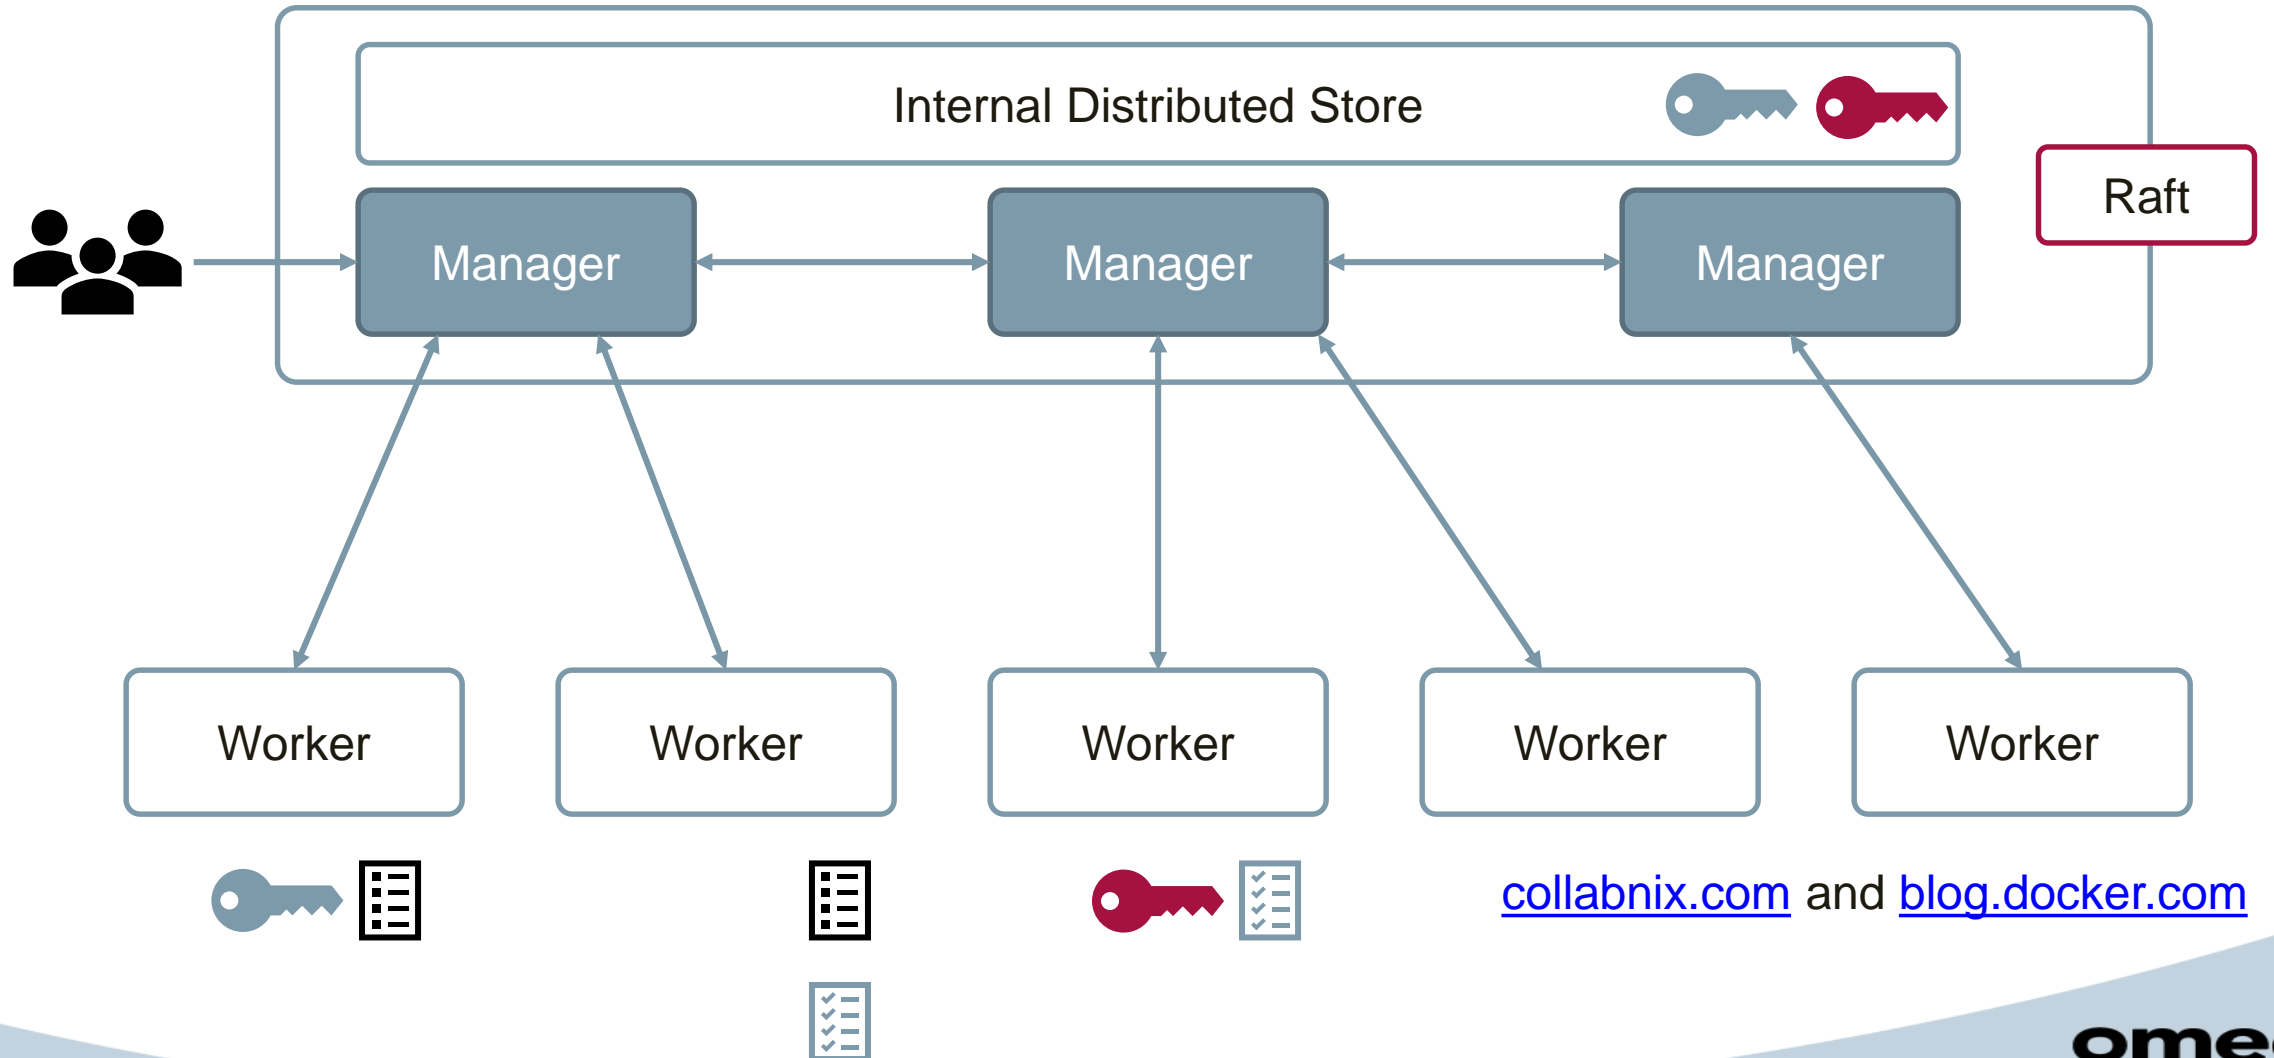
- Encrypted Raft logs
- Encrypted control and management plane traffic with mutual TLS
- HA Scheduling
- (optional) Encryption for application data plane traffic



# Docker Secrets

- Encrypted
  - In transit
  - At rest
- Centrally managed
- Mounted in an in-memory filesystem
  - /run/secrets/<secret\_name>

# Cluster



```
$ docker secret create mySecret secretInFile
```

```
$ docker service create --secret mySecret ...
```

```
$ docker create secrets site.key site.key  
$ docker create secrets site.crt site.crt
```

```
server {  
    listen          443 ssl;  
    server_name     localhost;  
    ssl_certificate  /run/secrets/site.crt;  
    ssl_certificate_key /run/secrets/site.key;  
  
    location / {  
        root       /usr/share/nginx/html;  
        index      index.html index.htm;  
    }  
}
```

```
$ docker service create \  
  --name nginx \  
  --secret site.key \  
  --secret site.crt \  
  --config source=site.conf, \  
            target=/etc/nginx/conf.d/site.conf, \  
            mode=0440 \  
  --publish published=3000,target=443 \  
nginx:latest
```

# Resources

- Documentation: <https://docs.docker.com>
- Docker on Youtube:  
<https://www.youtube.com/channel/UC76AVf2JkrwjxNKMUPpscHQ>
- Hands-on Learning: <https://training.play-with-docker.com>